



0. Course Details & Introduction

COMP 3270
Artificial Intelligence

Dirk Schnieders

0. Course Details & Introduction - Outline

- Course Details

- Prerequisite
- You will not learn ...
- Staff
- Syllabus
 - Topics
 - Schedule
- Assessment
- Plagiarism
- Slides
- Zoom
- Discussion Forum
- Textbook

- Introduction

- Intelligence vs AI
- Goal of AI
- Rational Agent
- Machine Learning
- History of AI
 - Turing Award Winners
 - Milestones
- The State of the Art
- Risks and Benefits of AI
 - Superhuman AI
 - The Gorilla Problem

Prerequisite

- COMP2119 Introduction to Data Structures and Algorithms (or equivalent)
 - Mathematics foundation
 - Growth of functions, recurrences, time and space complexity
 - Basic data structures and algorithms
 - Array, linked list, stack, queue, trees, hashing, graphs, sorting and searching algorithms
 - Advanced data structures and algorithms
 - Balanced tree, heap, priority queue, sorting in linear time

You will not learn ...

- Machine Learning: Perceptron, SVM, Linear Regression, ...
 - Consider taking: comp3314 Machine Learning
- Deep Learning (CNN, RNN, GAN, Deep RL)
 - Consider taking: comp3340 Applied deep learning, and/or
 - Consider taking: comp3317 Computer vision
- To work on bigger applied projects
 - Consider taking: comp3414 Experiential learning on AI and robotics
- R language with a focus on statistical modeling
 - Consider taking: comp3354 Statistical learning
- Advanced NLP models: Take COMP3361

Staff

- Instructor: Dirk Schnieders
 - Email: sdirk@cs.hku.hk
 - Consultation hours: Tue 2:30 - 3:30 & by appointment
- TAs
 - Cao Yukang
 - Q1, A1, A2, Forum admin during time: [start of course, Q1]
 - Email: ykcao@cs.hku.hk
 - Consultation hours
 - Thursday 2:30-3:30, Room 411 Chow Yei Ching Building
 - Guo Qiushan
 - Q2, A3, A4, Forum admin during time: [Q1, final exam]
 - Email: u3007956@connect.hku.hk
 - Consultation hours: Thu 2:30-3:30, Room P501 - Graduate Hall

Syllabus - Topics*

*subject to change

- Part I: Search and Planning
 - Search (Uninformed, Informed, Local, CSPs, Adversarial)
 - MDPs (Value Iteration, Policy Iteration)
 - RL (Temporal Difference Learning, Q Learning)
- Part II: Probabilistic Reasoning
 - MMs (Probability Review, Markov Chain, Mini-Forward Algorithm)
 - HMMs (Forward Algorithm, Particle Filtering)
 - BNs (Inference, Sampling)
- Part III: Intro to Machine Learning
 - NLP or another ML application (send me an email with topic requests)

Week	Tue (1h)	Topic*	Fri (2h)	Topic*
1 (1 Sep - 2 Sep)			2 Sep	Course Details & Introduction
2 (5 Sep - 9 Sep)	6 Sep	Uninformed Search	9 Sep	Informed Search
3 (12 Sep - 16 Sep)	13 Sep	Local Search	16 Sep	CSPs
4 (19 Sep - 23 Sep)	20 Sep	Adversarial Search	23 Sep	Adversarial Search
5 (26 Sep - 30 Sep)	27 Sep	Markov Decision Processes	30 Sep	Markov Decision Processes
6 (3 Oct - 7 Oct)	4 Oct	no class	7 Oct	Reinforcement Learning
7 (10 Oct - 14 Oct)	11 Oct	no class	14 Oct	no class
8 (17 Oct - 21 Oct)	18 Oct	Quiz 1	21 Oct	Reinforcement Learning
9 (24 Oct - 28 Oct)	25 Oct	Markov Models	28 Oct	Markov Models
10 (31 Oct - 4 Nov)	1 Nov	Hidden Markov Models	4 Nov	Hidden Markov Models
11 (7 Nov - 11 Nov)	8 Nov	Hidden Markov Models	11 Nov	Bayes Nets
12 (14 Nov - 18 Nov)	15 Nov	Bayes Nets	18 Nov	Bayes Nets
13 (21 Nov - 25 Nov)	22 Nov	NLP	25 Nov	NLP
14 (28 Nov - 30 Nov)	29 Nov	Quiz 2		

Assessment - Weight

- Exam: 50%
- Quiz 1: 8%
- Quiz 2: 8%
- Assignment 1: 10%
- Assignment 2: 8%
- Assignment 3: 8%
- Assignment 4: 8%

$\text{mark} = Q1 + Q2 + A1 + A2 + A3 + A4 + \text{exam}$

$\text{grade} = m(\text{mark})$

the mapping function $m()$ will not be published

Assessment - Quiz 1 & 2

- Written, closed book, calculator allowed
- Plagiarism cases will receive 0 marks and referred to program director
- 0 marks unless sick leave certificate (on day of test) or official leave from university
- Quiz 1: Scope: Everything covered in course
- Quiz 2: Scope: Everything covered in course since Quiz 1

Assessment - Assignments

- 4 Programming Assignments
 - Python 3, Suggestion: Install [Anaconda 3](#)
 - ~2 weeks time, no deadline extension
 - Late submission, after the deadline, will result in 0 marks
 - Submission of wrong files will result in 0 marks
 - Please check very carefully
 - Release - Deadline (subject to change)
 - A1: 13 Sep - 27 Sep
 - A2: 27 Sep - 11 Oct
 - A3: 25 Oct - 8 Nov
 - A4: 11 Nov - 25 Nov

Assessment - Final Exam

- 2 hours
- Details to be announced

Assessment - Results

- You must check for the correctness of your in-class assessment on Moodle
 - If you find any problem, send an email to our TA immediately

Plagiarism

- What is Plagiarism ?
 - <https://tl.hku.hk/plagiarism/>
- First attempt
 - Written warning
 - Zero marks
 - Referred to program director
- Second attempt
 - Referred to University Disciplinary Committee
 - Published reprimand
 - Suspension of study
 - Expulsion

Plagiarism

- We will treat your submission as plagiarism if we have sufficient evidence
 - Source will be reported and will also receive 0 marks
- You will receive a formal warning by the department in writing



Prevent Plagiarism

- Never copy/paste anything **without citing**
 - If you copy existing work you must clearly state and cite appropriately
 - Solutions to assignments may be found on the internet
 - To prevent plagiarism, don't search for them
- Never look at existing solutions
 - Work on the problem yourself
- Never ask your friend to see his/her solution
 - Note that the source will also be punished with 0 marks

Slides

- Available on Moodle on the day before the lecture
- Some slides include comments/solutions that I use during the lecture
- Use the provided materials responsibly
 - Don't share with others
 - Do not upload to internet
 - I do not hold copyright for most materials

Recording


- All lectures will be held face2face only
- Recordings might be uploaded
 - Do not rely on the recording, there could be a technical issue or I could forget to switch it on
 - Please remind me to switch on the recording

Discussion Forum

- If you have questions about course materials, please **use the forum**
 - We will try to answer within 48 hours
 - Please help to answer questions by others
- If you contact us by email please use your university email and include your university number in the email

Textbook

- Artificial Intelligence: A Modern Approach (4th Edition)
 - <http://aima.cs.berkeley.edu/>



Introduction

Intelligence

- We call ourselves Homo sapiens
 - Latin: Wise Man
 - Intelligence is important to us
- For thousands of years, we have tried to understand how we humans think
- Intelligence is most widely studied in humans, but has also been observed in animals and in plants

Artificial Intelligence

- AI goes further than just understanding intelligence
 - It attempts to build intelligent entities
- Is AI science, or is it engineering?
 - AI's science goal
 - To understand the principles and mechanism that account for intelligent action
 - AI's engineering goal
 - To design intelligent artifacts that can survive and operate in the physical world and solve problems of considerable scientific difficulty at high levels of competence


Goal of AI

- Should we create machines that ...
 - A. Think Humanly
 - B. Act Humanly
 - C. Think Rationally
 - D. Act Rationally





A. Think Humanly

- Researchers find out how we think by
 - Introspection
 - Psychological experiments
 - Brain imaging
- Cognitive science constructs precise and testable theories of the human mind
 - E.g., express a theory as a computer program and compare input-output behaviors to a human
 - If there is a match, some of the programs mechanism could also be operating in humans



A. Think Humanly

- The human brain is one of the great mysteries of science
 - How does our brain process information?
- The brain consists of nerve cells (aka neurons) and the collection of these simple cells leads to thought, action and consciousness
- The recent development of functional magnetic resonance imaging (fMRI) provides neuroscientists with details of brain activities



A. Think Humanly

- Brains and digital computers have somewhat different properties
- A crude comparison of the raw computational resources

	Supercomputer	Personal Computer	Human Brain
Computational units	10^6 GPUs + CPUs	8 CPU cores	10^6 columns
	10^{15} transistors	10^{10} transistors	10^{11} neurons
Storage units	10^{16} bytes RAM	10^{10} bytes RAM	10^{11} neurons
	10^{17} bytes disk	10^{12} bytes disk	10^{14} synapses
Cycle time	10^{-9} sec	10^{-9} sec	10^{-3} sec
Operations/sec	10^{18}	10^{10}	10^{17}

- Would we be able to achieve the brain's level of intelligence with a computer of unlimited capacity?

Should we create machines that ...

- A. Think Humanly
- B. Act Humanly
- C. Think Rationally
- D. Act Rationally


B. Act Humanly



- The Turing Test (aka Imitation Game) was designed to provide a definition of intelligence
- A computer passes the test if a human interrogator, after posing some questions, cannot tell whether the response come from a human or a computer



Human
interrogator



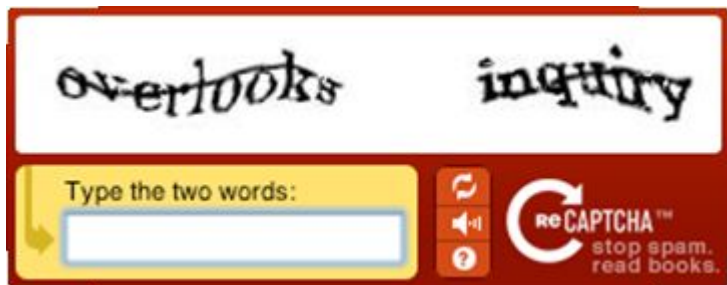
B. Act Humanly

- The underlying principles of intelligence are more important than to duplicate an exemplar
- Consider another field: Artificial Flight
 - The Wright brothers succeeded because they stopped imitating birds and started using wind tunnels and learn about aerodynamics
 - It was not their goal to make “machines that fly so exactly like pigeons that they can fool even other pigeons”



B. Act Humanly

- A reverse Turing test is a Turing test in which the objective or roles between computers and humans have been reversed. Interrogator is a computer. Interrogatee is a human.
 - Example: CAPTCHA
 - It is program that can generate and grade tests that
 - most humans can pass, but
 - current computer programs cannot pass



C. Think Rationally

- What are the laws that guide and underlie our thinking?
- Greek schools developed various forms of logic
 - Notation and rules of derivation for thoughts
 - Example: Socrates is a man; all men are mortal; therefore, Socrates is mortal
- By 1965, programs existed that could (in principle) solve any solvable problem described in logic notation
- Problems with this approach
 - How to take informal knowledge and state it in formal terms?
How about uncertainty?

D. Act Rationally

- Act so as to achieve the best outcome or, when there is uncertainty, the best expected outcome
- Advantages over the other approaches
 - More general than the laws of thought approach because correct inference is just one of several possible mechanisms for achieving rationality
 - The standard of rationality is mathematically well defined. Human behavior, on the other hand, is well adapted only for one specific environment
- In this course we will focus on the general principles of rational agents and how to build them
 - An agent is something that perceives and acts

Rational Agent

- This course is about designing rational agents
- For any given class of environments and tasks, we seek the agent (or class of agents) with the best performance
 - Computational limitations make perfect rationality unachievable
 - Design best program for given machine resources

Machine Learning

- An agent is learning if it improves its performance on future tasks
- Why would we want an agent to learn? If the design of an agent can be improved, why not design the agent with that improvement to begin with?
 - Cannot anticipate all possible situations
 - Cannot anticipate all changes over time
 - Don't know how to program some solutions

History of AI



History of AI - Turing Award Winners

- Marvin Minsky (1969)
- John McCarthy (1971)
- Edward Feigenbaum and Raj Reddy (1994)
- Judea Pearl (2011)
- Yoshua Bengio, Geoffrey Hinton, and Yann LeCun (2019)

History of AI - Milestones

- Inception (1943 - 1956)
- Early Enthusiasm (1952 - 1969)
- A dose of reality (1966 - 1973)
- Expert systems (1969 - 1986)
- Return of NN (1986 - present)
- Probabilistic reasoning (1987 - present)
- Big data (2001 - present)
- Deep Learning (2011 - present)

The State of the Art

- Publications
 - AI papers increased 20 fold between 2010 to 2019 to 20,000 a year
- Conferences
 - Attendance of [NeurIPS](#) increased 800% since 2012 to 13,500
- Industry
 - AI startups in the US increased 20 fold to over 800 from 2010 to 2019
- Internationalization
 - China publishes more AI papers per year then US and about as many as Europe
 - In citation weighted impact, US is ahead by 50% vs. China

The State of the Art

- Vision
 - Error rates for object detection improved from 28% to less than 2%
- Speed
 - Training time for image recognition dropped by a factor of 100 in last 2 years
 - Amount of computing power used in top AI applications is doubling every few month
- Humans vs. AI
 - AI is better in chess, go, poker, pac-man, jeopardy!, object detection, speech recognition in limited domain, chinese-to-english in restricted domain, Quake III, Dota 2, StarCraft II, many Atari games, Skin cancer detection, prostate cancer detection, protein folding, ...

Benefits of AI

First solve AI, then use AI to solve everything else.

Demis Hassabis, Google DeepMind



Risks of AI

- Lethal autonomous weapons
- Surveillance
- Biased decision making
- Impact on employment
- Safety-critical applications
- Cybersecurity

Risks of AI - Superhuman AI

- Most experts agree that we will eventually be able to create a superhuman AI
 - An intelligence that far surpasses human ability

Risks of AI - The Gorilla Problem

- About seven million years ago, a now-extinct primate evolved
 - one branch led to gorillas
 - another to humans
- Today the gorillas are probably not too happy about the human branch
 - They have no control over their future



Risks of AI - The Gorilla Problem

- If the gorilla problem is the result of developing AI then we should stop working on it
- If superhuman AI were a black box from outer space, we should be careful in opening the box
 - But it is not, **we** design the AI systems
 - If AI does end up taking control, it would be a design failure
- We need to understand the source of potential failure
 - Philosophical foundations of AI
 - Maybe the most important area of AI research

Reference

- Textbook: Chapter 1